Dear Hiring Manager:

I recently completed an M.S. in Computer Science at UCLA, and I'd love to join Kaiser as a Data Analyst. I enjoy poking and prodding at data to see what can be extracted as well as learning new things. In addition, during my eight years at the L.A. Times, I gained extensive experience analyzing and using consumer data.

As a Senior Editor for Digital at the L.A. Times, one of my roles was to create traffic reports for sections across the paper, including local, entertainment, features and sports. This involved communicating with departments to discover their needs and then creating automated reports through Adobe Analytics (Omniture), a Web-based tool for tracking user activity. Typically, I created daily and weekly reports for sections that showed pageviews and visits to content. In addition to automated reports, I also created custom reports that involved restructuring data. For example, I compiled a list of the most- and least-viewed articles for each section over a six-month period.

In addition to creating reports, I also used consumer data to guide content creation and placement. I used reader viewing habits to determine when to place stories, where to place them and when to remove them. For example, the most visible locations and timeslots were saved for the most important stories or those most likely to receive the widest viewership. Using insights derived from traffic reports, I helped editors and writers understand what online readers were interested in viewing. For example, in the travel section, my observation that readers liked content about obscure and unique places lead to increased articles on this topic.

In terms of data science experience, I've worked on several academic projects involving data analytics and machine learning. These projects exposed me to data cleaning, feature engineering, model validation, different machine learning algorithms as well as distributed computing frameworks such as Spark and Hadoop.

For a Big Data class, I worked in a group of three to perform cross-device entity resolution using anonymized user Web browsing histories (gigabytes worth of it). I took the lead on implementing the project, establishing the approach and dividing work among team members. Our entity resolution approach used regression to perform binary classification. Using the browsing histories, we extracted features such as URL title n-grams and access times to assess similarities between users. Our implementation was done using Python and the scikit-learn machine learning library.

For my M.S. degree project, I built on skills gained through a Health Analytics class to develop a stress detection system. The system uses an Apple Watch to collect data used for modeling and detecting stress. During the supervised learning period, the watch asks the user to input her or his stress level and then records other measurements for modeling, including heart rate, resting heart rate, step count and time of day. After a sufficient number of training samples have been collected, a model is created and used to detect stress. Stress level is visualized on the watch. A history of stress levels is graphed on a companion iOS app.

Thanks for getting to know me a little, and I'd love to speak to you about this position.

Best,
Jason La